




University of Bari — Department of Computer Science — Department of Physics  
LACAM Machine Learning Lab — National Institute for Nuclear Physics

# Fast and Accurate Density Estimation with Extremely Randomized Cutset Networks

Nicola Di Mauro  **Antonio Vergari**  Teresa M.A. Basile Floriana Esposito

 = both authors contributed equally

19th September - **ECML-PKDD 2017** - Skopje, Macedonia

# Outline

- Density Estimation
- Tractable Probabilistic Models
- Cutset Networks
- XC Nets
- Experiments
- Conclusions
- *References*

# Density Estimation

**Density estimation** is the unsupervised task of learning an estimator for the joint probability distribution  $p(\mathbf{X})$  from a set of i.i.d. samples  $\mathcal{D} = \{\mathbf{x}^i\}_{i=1}^m$  over random variables (RVs)  $\mathbf{X} = \{X_1, \dots, X_n\}$



Given such an estimator, one uses it to *answer probabilistic queries* about configurations on  $\mathbf{X}$ , i.e. to do **inference**. E.g., classification can be performed by Most Probable Explanation (MPE) inference:  $y^* = \operatorname{argmax}_{y \sim Y} p(y|\mathbf{X})$ .

The main challenge in density estimation is balancing:

- ▶ the **representation expressiveness** of the model to learn
- ▶ the **cost of learning** such a model
- ▶ and the **cost of performing inference** on it

# Tractable Probabilistic Models (TPMs)

Classical Probabilistic Graphical Models like **Bayesian Networks (BNs)** and **Markov Networks (MNs)** are highly expressive but exact inference is generally **NP-hard** with them [Roth 1996].




**Tractable Probabilistic Models (TPMs)** on the other hand, are density estimators for which some kind of **exact inference is tractable**, i.e. *polynomial* in the number of RVs, i.e.,  $n$ , or their domains

→ learning may still be hard to scale on high-dimensional data

# Product of Bernoullis (PoBs)

A not so much expressive TPM, assuming all RVs to be independent:




$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i)$$

Learning a PoB has linear time complexity  $O(nm)$

Complete evidence inference is linear  $O(n)$

# Chow-Liu Trees (CLTrees)

A *directed tree-structured model* [Meilă and Jordan 2000] over  $\mathbf{X}$  is a BN in which each node  $X_i \in \mathbf{X}$  has at most one parent,  $\text{Pa}_{X_i}$ .




$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | \text{Pa}_{x_i})$$

Complete evidence inference is still linear for CLtrees:  $O(n)$

But learning now takes quadratic time  $O(n^2(m + \log n))$


# Cutset Networks (CNets)



A **Cutset Network** (CNet)  $\mathcal{C}$  is a TPM represented via a **weighted probabilistic model** tree over  $\mathbf{X}$  recursively defined as:

1. a TPM  $\mathcal{M}$ , with  $\text{scope}(\mathcal{M}) = \mathbf{X}$
2. a weighted disjunction (OR node) of two CNets  $\mathcal{C}_0$  and  $\mathcal{C}_1$  conditioned on RV  $X_i \in \mathbf{X}$ , with weights  $w_i^0$  and  $w_i^1$  s.t.  $w_i^0 + w_i^1 = 1$ , where  $\text{scope}(\mathcal{C}_0) = \text{scope}(\mathcal{C}_1) = \mathbf{X}_{\setminus i}$

# CNets I




A CNet  $\mathcal{C}$  defines the following joint distribution:

$$p(\mathbf{x}) = p_l(\mathbf{x}_{\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{\text{scope}(\mathcal{M}_l)})$$




# CNets II



A CNet  $\mathcal{C}$  acts as a **deterministic mixture of experts** in which the OR tree acts as the **gating function**

$$p(\mathbf{x}) = p_l(\mathbf{x}_{|\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{|\text{scope}(\mathcal{M}_l)})$$

# CNets III



and in which leaf models  $\mathcal{M}_l$  play the role of **local experts**

$$p(\mathbf{x}) = p_l(\mathbf{x}_{|\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{|\text{scope}(\mathcal{M}_l)})$$

complete evidence inference is still linear  $O(n)$  !

# Learning C Nets I

**Top-down greedy** CNet learners can be unified in single template, LearnCNet:

---

LearnCNet( $\mathcal{D}, \mathbf{X}, \delta, \sigma$ )

- 1: **Input:** a dataset  $\mathcal{D}$  over RVs  $\mathbf{X}$ ;  $\delta$  min number samples;  $\sigma$  min number features
  - 2: **Output:** a CNet  $\mathcal{C}$  encoding  $p_{\mathcal{C}}(\mathbf{X})$  learned from  $\mathcal{D}$
  - 3: **if**  $|\mathcal{D}| > \delta$  **and**  $|\mathbf{X}| > \sigma$  **then**
  - 4:    $X_i \leftarrow \text{select}(\mathcal{D}, \mathbf{X})$
  - 5:    $\mathcal{D}_0 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 0\}, \mathcal{D}_1 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 1\}$
  - 6:    $w_0 \leftarrow |\mathcal{D}_0|/|\mathcal{D}|, w_1 \leftarrow |\mathcal{D}_1|/|\mathcal{D}|$
  - 7:    $\mathcal{C} \leftarrow w_0 \cdot \text{LearnCNet}(\mathcal{D}_0, \mathbf{X}_{\setminus i}, \delta, \sigma) + w_1 \cdot \text{LearnCNet}(\mathcal{D}_1, \mathbf{X}_{\setminus i}, \delta, \sigma)$
  - 8: **else**
  - 9:    $\mathcal{C} \leftarrow \text{learnLeafDistribution}(\mathcal{D}, \mathbf{X})$  ▷ Grow a leaf TPM
  - 10: **return**  $\mathcal{C}$
- 

in which for the base case, if no conditioning is possible, *a leaf distribution is estimated*

# Learning C Nets II

**Top-down greedy** CNet learners can be unified in single template, LearnCNet:

---

LearnCNet( $\mathcal{D}, \mathbf{X}, \alpha, \delta, \sigma$ )

- 1: **Input:** a dataset  $\mathcal{D}$  over RVs  $\mathbf{X}$ ;  $\delta$  min number samples;  $\sigma$  min number features
  - 2: **Output:** a CNet  $\mathcal{C}$  encoding  $p_{\mathcal{C}}(\mathbf{X})$  learned from  $\mathcal{D}$
  - 3: **if**  $|\mathcal{D}| > \delta$  **and**  $|\mathbf{X}| > \sigma$  **then**
  - 4:    $X_i \leftarrow \text{select}(\mathcal{D}, \mathbf{X})$  ▷ select the RV to condition on
  - 5:    $\mathcal{D}_0 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 0\}, \mathcal{D}_1 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 1\}$
  - 6:    $w_0 \leftarrow |\mathcal{D}_0|/|\mathcal{D}|, w_1 \leftarrow |\mathcal{D}_1|/|\mathcal{D}|$
  - 7:    $\mathcal{C} \leftarrow w_0 \cdot \text{LearnCNet}(\mathcal{D}_0, \mathbf{X}_{\setminus i}, \delta, \sigma) + w_1 \cdot \text{LearnCNet}(\mathcal{D}_1, \mathbf{X}_{\setminus i}, \delta, \sigma)$
  - 8: **else**
  - 9:    $\mathcal{C} \leftarrow \text{learnLeafDistribution}(\mathcal{D}, \mathbf{X})$
  - 10: **return**  $\mathcal{C}$
- 

or selecting a RV to condition on is performed, splitting the dataset and recursing

# Learning C Nets III

Different implementations of select lead to different time complexities:

**entCNet** choosing  $X_i$  to lower approximate average joint entropy [Rahman, Kothalkar, and Gogate 2014]

$$\rightarrow O(mn^2)$$

**dCSN** choosing  $X_i$  in a principled way improving likelihood [Di Mauro, Vergari, and Esposito 2015]

$$\rightarrow O(n^3(m + \log n))$$

Quadratic and cubic times for *each* RV selection do not scale on high dimensional data  $\rightarrow$  We aim at drastically reducing it!

# XCNets I

**XCNets (Extremely Randomized CNets)** are CNets built by LearnCNet when **select** chooses one RV *completely at random*.

select time complexity →  $O(1)$ !

Advantages of a **single** XCNet over a classically learned CNet:

- ▶ extremely fast to learn
- ▶ only slightly less accurate as density estimators
- ▶ almost as good at generating samples
- ▶ less prone to overfitting

When plugged into **ensembles** they outperform state-of-the-art density estimators in a fraction of the time!

# XCNets II

Why a single XCNet is *not much worse* than a CNet?

Because of the mixture of experts interpretations, a path  $p = p_{(1)}p_{(2)} \cdots p_{(k)}$  connects the root to a single leaf  $\mathcal{M}_l$  after observing  $x_1x_2 \cdots x_k$ ,

$p_l(\mathbf{x}_{|\text{sc}(\mathcal{C}) \setminus \text{sc}(\mathcal{M}_l)})$  is decomposed by the chain rule across path  $p$

→ shuffling  $X_{p_{(1)}}, \dots, X_{p_{(k)}}$  does not influence inference!

Focus on learning **accurate local experts** (leaves) than the gating function!

Sampling is less affected:




CNet



Train



XCNet





Train

# XCNets III

Single C Nets learned with LearnCNet are prone to overfitting.

***Randomizing the gating function*** in XC Nets alleviate this issue.



Learning curves of C Nets and XC Nets (negative log-likelihoods) on KddCup2k and EachMovie The latter overfits much later than the former.

Moreover, it helps differentiating the local leaf experts



# Ensembles of XC Nets

Ensembles of XC Nets do not require to additionally diversify each component (e.g. no bootstrapping required).

Learning up to **500 components** of XC Nets is still **faster than learning 40** models of other variants:

**CNet<sub>bag</sub>** bagging entCSN [Rahman and Gogate 2016]

**CNet<sub>boost</sub>** boosting entCSN [Rahman and Gogate 2016]

**dCSN<sup>k</sup>** bagging dCSN [Di Mauro, Vergari, and Basile 2015;  
Di Mauro, Vergari, and Esposito 2015]

# Experiments

We validate the following research questions:

- (Q1) how does a single XCNet **compare** to the optimal one learned by dCSN?
- (Q2) how **accurate** are ensembles of XC Nets compared to state-of-the-art density estimators?
- (Q3) how **much time** do actually XC Nets save in practice?

We employ the **20 standard benchmark datasets** for density estimation [*Haaren and Davis 2012; Vergari, Mauro, and Esposito 2015*]

We compare to single and ensembles of C Nets plus other state-of-the-art TPMs like Sum-Product Networks (ID-SPN) and Markov Networks learned with ACs (ACMN) and even untractable Bayesian Networks (BN).

# (Q1) Single Model Comparison

dataset	entCNet	dCSN	XCNet	dCSN <sub>PoB</sub>	XCNet <sub>PoB</sub>
NLTCS	-6.06	<b>-6.03</b>	$6.06 \pm 0.01$	-6.09	$-6.17 \pm 0.05$
MSNBC	<b>-6.05</b>	<b>-6.05</b>	$-6.09 \pm 0.02$	-6.05	$-6.18 \pm 0.03$
KDDCup2k	-	<b>-2.18</b>	$-2.19 \pm 0.01$	-2.19	$-2.21 \pm 0.01$
Plants	<b>-13.25</b>	<b>-13.25</b>	$-13.43 \pm 0.07$	-14.89	$-15.66 \pm 0.22$
Audio	<b>-42.05</b>	-42.10	$-42.66 \pm 0.14$	-42.95	$-44.02 \pm 0.22$
Jester	-55.56	<b>-55.40</b>	$-56.10 \pm 0.19$	-56.23	$-57.39 \pm 0.15$
Netflix	<b>-58.71</b>	<b>-58.71</b>	$-59.21 \pm 0.06$	-60.20	$-61.40 \pm 0.25$
Accidents	-30.69	<b>-29.84</b>	$-31.58 \pm 0.24$	-36.24	$-40.22 \pm 0.46$
Retail	<b>-10.94</b>	-11.24	$-11.44 \pm 0.09$	-11.06	$-11.19 \pm 0.04$
Pumsb-star	-24.42	<b>-23.91</b>	$-25.55 \pm 0.34$	-32.11	$-39.91 \pm 2.48$
DNA	-87.59	<b>-87.31</b>	$-87.67 \pm 0.00$	-98.83	$-99.84 \pm 0.05$
Kosarek	<b>-11.04</b>	-11.20	$-11.70 \pm 0.13$	-11.38	$-11.80 \pm 0.07$
MSWeb	<b>-10.07</b>	-10.10	$-10.47 \pm 0.10$	-10.19	$-10.43 \pm 0.07$
Book	<b>-37.35</b>	-38.93	$-42.36 \pm 0.28$	-38.21	$-39.47 \pm 0.33$
EachMovie	-58.37	<b>-58.06</b>	$-60.71 \pm 0.89$	-59.70	$-62.58 \pm 0.38$
WebKB	-162.17	<b>-161.92</b>	$-167.45 \pm 1.59$	-168.7	$-174.78 \pm 0.81$
Reuters-52	<b>-88.55</b>	-88.65	$-99.52 \pm 1.93$	-90.51	$-100.25 \pm 0.57$
20NewsG	-	<b>-161.72</b>	$-172.6 \pm 1.40$	-162.25	$-167.39 \pm 0.74$
BBC	-263.08	<b>-261.79</b>	$-261.79 \pm 0.00$	-264.56	$-274.83 \pm 1.15$
Ad	-16.92	<b>-16.34</b>	$-18.70 \pm 1.44$	-36.44	$-36.94 \pm 1.41$

Table 1.

Average test log-likelihoods for entCNet, dCSN, XCNet and their PoB variants dCSN<sub>PoB</sub> and XCNet<sub>PoB</sub>. For randomized models, mean and standard deviation over 10 runs are reported.

## (Q2) Ensemble Model Comparison

new *state-of-the-art* scores on 11/20 datasets...

dataset	CNet <sub>40</sub> <sub>bag</sub>	CNet <sub>40</sub> <sub>boost</sub>	dCSN <sub>40</sub>	XCNet <sub>40</sub> <sub>p08</sub>	XCNet <sub>40</sub>	XCNet <sub>500</sub>	ID-SPN	ACMN	WM
NLTCS	<b>-6.00</b>	-6.01	<b>-6.00</b>	-6.01	<b>-6.00</b>	<b>-5.99</b>	-6.02	-6.00	-6.02
MSNBC	-6.08	-6.15	<b>-6.05</b>	-6.11	-6.06	-6.06	<b>-6.04</b>	<b>-6.04</b>	<b>-6.04</b>
KDDCup2k	-2.14	-2.15	-2.15	<b>-2.13</b>	<b>-2.13</b>	<b>-2.13</b>	<b>-2.13</b>	-2.17	-2.16
Plants	-12.32	-12.67	-12.59	-13.09	<b>-11.99</b>	<b>-11.84</b>	-12.54	-12.80	-12.65
Audio	-40.09	-39.84	-40.19	-40.30	<b>-39.77</b>	<b>-39.39</b>	-39.79	-40.32	-40.50
Jester	-52.88	-52.82	-52.99	-53.64	<b>-52.65</b>	<b>-52.21</b>	-52.86	-53.31	-53.85
Netflix	-56.55	-56.44	-56.69	-57.64	<b>-56.38</b>	<b>-55.93</b>	-56.36	-57.22	-57.03
Accidents	-29.88	-29.45	<b>-29.27</b>	-36.92	-29.31	-29.10	-26.98	-27.11	<b>-26.32</b>
Retail	-10.84	<b>-10.81</b>	-11.17	-10.88	-10.93	-10.91	-10.85	-10.88	-10.87
Pumsb-star	-23.98	-23.46	-23.78	-32.91	<b>-23.44</b>	-23.31	-22.41	-23.55	<b>-21.72</b>
DNA	<b>-81.07</b>	-85.67	-85.95	-98.28	-84.96	-84.17	-81.21	<b>-80.03</b>	-80.65
Kosarek	-10.74	<b>-10.60</b>	-10.97	-10.91	-10.72	-10.66	<b>-10.60</b>	-10.84	-10.83
MSWeb	-9.77	-9.74	-9.93	-9.83	<b>-9.66</b>	<b>-9.62</b>	-9.73	-9.77	-9.70
Book	-35.55	<b>-34.46</b>	-37.38	-34.77	-36.35	-35.45	<b>-34.14</b>	-35.56	-36.41
EachMovie	-53.00	<b>-51.53</b>	-54.14	-51.66	-51.72	<b>-50.34</b>	-51.51	-55.80	-54.37
WebKB	-153.12	<b>-152.53</b>	-155.47	-155.83	-153.01	<b>-149.20</b>	-151.84	-159.13	-157.43
Reuters-52	-83.71	<b>-83.69</b>	-86.19	-85.16	-84.05	<b>-81.87</b>	-83.35	-90.23	-87.55
20NewsG	-156.09	-153.12	-156.46	<b>-152.21</b>	-153.89	<b>-151.02</b>	-151.47	-161.13	-158.95
BBC	<b>-237.42</b>	-247.01	-248.84	-251.31	-238.47	<b>-229.21</b>	-248.93	-257.10	-257.86
Ad	-15.28	-14.36	-15.55	-26.25	<b>-14.20</b>	<b>-14.00</b>	-19.05	-16.53	-18.35

## (Q3) Learning Times

...in a ***fraction of the time*** required by other competitors

dataset	dCSN	XCNet	dCSN <sub>PoB</sub>	XCNet <sub>PoB</sub>	dCSN <sup>40</sup>	XCNet <sub>PoB</sub> <sup>40</sup>	XCNet <sup>40</sup>	XCNet <sup>500</sup>	ID-SPN
NLTCS	7	0.2	0.1	0.01	10	0.2	0.01	3	310
MSNBC	12	0.3	0.7	0.01	499	13.1	13	155	46266
KDDCup2k	112	0.5	12.0	0.32	4126	21.2	16	247	32067
Plants	15	0.3	45.5	0.22	325	1.0	6	77	18833
Audio	58	0.3	74.8	0.48	980	0.8	6	136	21009
Jester	50	0.2	95.6	0.26	989	0.3	4	83	10412
Netflix	75	0.2	2.8	0.02	1546	0.4	9	118	30294
Accidents	54	0.2	153.7	0.04	996	0.7	11	138	15472
Retail	263	0.8	5.8	0.01	3780	3.2	13	164	4041
Pumsb-star	118	0.6	26.2	0.02	2260	0.8	23	290	20952
DNA	30	0.1	4.4	0.01	224	0.06	3	40	3040
Kosarek	588	2.4	41.2	0.01	10033	10.8	43	524	17799
MSWeb	1215	7.2	7.4	0.01	17123	13.2	129	1592	19682
Book	9235	9.7	113.0	0.04	155634	1.9	316	3476	61248
EachMovie	1297	7.1	4.7	0.01	16962	1.1	127	2601	118782
WebKB	4997	11.0	238.0	0.03	18875	0.9	190	2237	45451
Reuters-52	9947	39.3	24.3	0.05	65498	2.7	414	8423	70863
20NewsG	16866	51.3	40.7	0.01	153908	4.4	506	9883	163256
BBC	21381	8.4	7.3	0.02	69572	0.4	256	4251	61471
Ad	5212	116.5	134.0	0.08	75694	4.2	2403	30538	87522

Table 3. Times (in seconds) taken to learn the best models on each dataset for dCSN, XCNet, dCSN<sub>PoB</sub>, XCNet<sub>PoB</sub>, their ensembles and ID-SPN.

# Conclusions

Due to their simplicity to implement, *fast learning times*, and *accurate inference performances*, XCNETs set the **new baseline to compare against** for density estimation with TPMs (and not!).

## Future works

Exploiting the mixture of experts interpretation to devise more expressive gating functions to still perform exact and fast inference

## Code

<https://github.com/nicoladimauro/cnet>

## Paper

<http://www.di.uniba.it/~ndm/pubs/ndm17ecml.pdf>

# References

- ⊕ Di Mauro, Nicola, Antonio Vergari, and Teresa M.A. Basile (2015). "Learning Bayesian Random Cutset Forests". In: *Proceedings of ISMIS*. Springer, pp. 122–132.
- ⊕ Di Mauro, Nicola, Antonio Vergari, and Floriana Esposito (2015). "Learning Accurate Cutset Networks by Exploiting Decomposability". In: *Proceedings of AIXIA*. Springer, pp. 221–232.
- ⊕ Haaren, Jan Van and Jesse Davis (2012). "Markov Network Structure Learning: A Randomized Feature Generation Approach". In: *Proceedings of the 26th Conference on Artificial Intelligence*. AAAI Press.
- ⊕ Meilä, Marina and Michael I. Jordan (2000). "Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1, pp. 1–48.
- ⊕ Rahman, Tahrira and Vibhav Gogate (2016). "Learning Ensembles of Cutset Networks". In: *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*. AAAI'16. Phoenix, Arizona: AAAI Press, pp. 3301–3307. URL: <http://dl.acm.org/citation.cfm?id=3016100.3016365>.
- ⊕ Rahman, Tahrira, Prasanna Kothalkar, and Vibhav Gogate (2014). "Cutset Networks: A Simple, Tractable, and Scalable Approach for Improving the Accuracy of Chow-Liu Trees". In: *Machine Learning and Knowledge Discovery in Databases*. Vol. 8725. LNCS. Springer, pp. 630–645.
- ⊕ Roth, Dan (1996). "On the hardness of approximate reasoning". In: *Artificial Intelligence* 82.1–2, pp. 273–302.
- ⊕ Vergari, Antonio, Nicola Di Mauro, and Floriana Esposito (2015). "Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning". In: *ECML-PKDD 2015*.

## Discuss