




University of Bari — Department of Computer Science — Department of Physics  
LACAM Machine Learning Lab — National Institute for Nuclear Physics

# Fast and Accurate Density Estimation with Extremely Randomized Cutset Networks

Nicola Di Mauro  **Antonio Vergari**  Teresa M.A. Basile Floriana Esposito

 = both authors contributed equally

19th September - **ECML-PKDD 2017** - Skopje, Macedonia

# Outline

- Density Estimation
- Tractable Probabilistic Models
- Cutset Networks
- XC Nets
- Experiments
- Conclusions
- *References*

# Density Estimation

**Density estimation** is the unsupervised task of learning an estimator for the joint probability distribution  $p(\mathbf{X})$  from a set of i.i.d. samples  $\mathcal{D} = \{\mathbf{x}^i\}_{i=1}^m$  over random variables (RVs)  $\mathbf{X} = \{X_1, \dots, X_n\}$



Given such an estimator, one uses it to *answers probabilistic queries* about configurations on  $\mathbf{X}$ , i.e. to do **inference**. E.g., classification can be performed by Most Probable Explanation (MPE) inference:  $y^* = \operatorname{argmax}_{y \sim Y} p(y|\mathbf{X})$ .

The main challenge in density estimation is balancing:

- ▶ the **representation expressiveness** of the model to learn
- ▶ the **cost of learning** such a model
- ▶ and the **cost of performing inference** on it.

# Tractable Probabilistic Models (TPMs)

Classical Probabilistic Graphical Models like **Bayesian Networks (BNs)** and **Markov Networks (MNs)** are highly expressive but exact inference is generally **NP-hard** with them [Roth 1996].




**Tractable Probabilistic Models (TPMs)** on the other hand, are density estimators for which some kind of **exact inference is tractable**, i.e. *polynomial* in the number of RVs, i.e.,  $n$ , or their domains.

→ learning may still be hard to scale on high-dimensional data

# Product of Bernoullis (PoBs)

A not so much expressive TPM, assuming all RVs to be independent:




$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i)$$

Learning a PoB has linear time complexity  $O(nm)$ .

Complete evidence inference is linear  $O(n)$ .

# Chow-Liu Trees (CLTrees)

A *directed tree-structured model* [Meilă and Jordan 2000] over  $\mathbf{X}$  is a BN in which each node  $X_i \in \mathbf{X}$  has at most one parent,  $\text{Pa}_{X_i}$ .




$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | \text{Pa}_{x_i})$$

Complete evidence inference is still linear for CLtrees:  $O(n)$ .

But learning now takes quadratic time  $O(n^2(m + \log n))$ .


# Cutset Networks (CNets)



A **Cutset Network** (CNet)  $\mathcal{C}$  is a TPM represented via a **weighted probabilistic model** tree over  $\mathbf{X}$  recursively defined as:

1. a TPM  $\mathcal{M}$ , with  $\text{scope}(\mathcal{M}) = \mathbf{X}$
2. a weighted disjunction (OR node) of two CNets  $\mathcal{C}_0$  and  $\mathcal{C}_1$  conditioned on RV  $X_i \in \mathbf{X}$ , with weights  $w_i^0$  and  $w_i^1$  s.t.  $w_i^0 + w_i^1 = 1$ , where  $\text{scope}(\mathcal{C}_0) = \text{scope}(\mathcal{C}_1) = \mathbf{X}_{\setminus i}$

# CNets I




A CNet  $\mathcal{C}$  defines the following joint distribution:

$$p(\mathbf{x}) = p_l(\mathbf{x}_{\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{\text{scope}(\mathcal{M}_l)})$$




# CNets II



A CNet  $\mathcal{C}$  acts as a **deterministic mixture of experts** in which the OR tree acts as the **gating function**

$$p(\mathbf{x}) = p_l(\mathbf{x}_{|\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{|\text{scope}(\mathcal{M}_l)})$$

# CNets III



and in which leaf models  $\mathcal{M}_l$  play the role of **local experts**

$$p(\mathbf{x}) = p_l(\mathbf{x}_{|\text{scope}(\mathcal{C}) \setminus \text{scope}(\mathcal{M}_l)}) p_{\mathcal{M}_l}(\mathbf{x}_{|\text{scope}(\mathcal{M}_l)})$$

complete evidence inference is still linear  $O(n)$  !

# Learning C Nets I

**Top-down greedy** CNet learners can be unified in single template, LearnCNet:

---

LearnCNet( $\mathcal{D}, \mathbf{X}, \alpha, \delta, \sigma$ )

- 1: **Input:** a dataset  $\mathcal{D}$  over RVs  $\mathbf{X}$ ;  $\alpha$ :  $\delta$  min number samples;  $\sigma$  min number features
- 2: **Output:** a CNet  $\mathcal{C}$  encoding  $p_{\mathcal{C}}(\mathbf{x})$  learned from  $\mathcal{D}$
- 3: **if**  $|\mathcal{D}| > \delta$  **and**  $|\mathbf{X}| > \sigma$  **then**
- 4:    $X_i \leftarrow \text{select}(\mathcal{D}, \mathbf{X}, \alpha)$
- 5:    $\mathcal{D}_0 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 0\}, \mathcal{D}_1 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 1\}$
- 6:    $w_0 \leftarrow |\mathcal{D}_0|/|\mathcal{D}|, w_1 \leftarrow |\mathcal{D}_1|/|\mathcal{D}|$
- 7:    $\mathcal{C} \leftarrow w_0 \cdot \text{LearnCNet}(\mathcal{D}_0, \mathbf{X}_{\setminus i}, \alpha, \delta, \sigma) + w_1 \cdot \text{LearnCNet}(\mathcal{D}_1, \mathbf{X}_{\setminus i}, \alpha, \delta, \sigma)$
- 8: **else**
- 9:    $\mathcal{C} \leftarrow \text{learnLeafDistribution}(\mathcal{D}, \mathbf{X}, \alpha)$  ▷ Grow a leaf TPM
- 10: **return**  $\mathcal{C}$

---

in which for the base case, if no conditioning is possible, *a leaf distribution is estimated*

# Learning C Nets II

**Top-down greedy** CNet learners can be unified in single template, LearnCNet:

---

LearnCNet( $\mathcal{D}, \mathbf{X}, \alpha, \delta, \sigma$ )

- 1: **Input:** a dataset  $\mathcal{D}$  over RVs  $\mathbf{X}$ ;  $\alpha$ :  $\delta$  min number samples;  $\sigma$  min number features
  - 2: **Output:** a CNet  $\mathcal{C}$  encoding  $p_{\mathcal{C}}(\mathbf{x})$  learned from  $\mathcal{D}$
  - 3: **if**  $|\mathcal{D}| > \delta$  **and**  $|\mathbf{X}| > \sigma$  **then**
  - 4:    $X_i \leftarrow \text{select}(\mathcal{D}, \mathbf{X}, \alpha)$  ▷ select the RV to condition on
  - 5:    $\mathcal{D}_0 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 0\}, \mathcal{D}_1 \leftarrow \{\xi \in \mathcal{D} : \xi[X_i] = 1\}$
  - 6:    $w_0 \leftarrow |\mathcal{D}_0|/|\mathcal{D}|, w_1 \leftarrow |\mathcal{D}_1|/|\mathcal{D}|$
  - 7:    $\mathcal{C} \leftarrow w_0 \cdot \text{LearnCNet}(\mathcal{D}_0, \mathbf{X}_{\setminus i}, \alpha, \delta, \sigma) + w_1 \cdot \text{LearnCNet}(\mathcal{D}_1, \mathbf{X}_{\setminus i}, \alpha, \delta, \sigma)$
  - 8: **else**
  - 9:    $\mathcal{C} \leftarrow \text{learnLeafDistribution}(\mathcal{D}, \mathbf{X}, \alpha)$
  - 10: **return**  $\mathcal{C}$
- 

or selecting a RV to condition on is performed, splitting the dataset and recursing

# Learning C Nets III

Different implementations of select lead to different time complexities:

**entCNet** choosing  $X_i$  to lower approximate average joint entropy [Rahman, Kothalkar, and Gogate 2014]

$$\rightarrow O(mn^2)$$

**dCSN** choosing  $X_i$  in a principled way improving likelihood [Di Mauro, Vergari, and Esposito 2015]

$$\rightarrow O(n^3(m + \log n))$$

Quadratic and cubic times for *each* RV selection do not scale on high dimensional data

We aim at drastically reducing it!

# XCNets I

**XCNets (Extremely Randomized CNets)** are CNets built by LearnCNet when select chooses one RV *completely at random*.

select time complexity →  $O(1)$ !

Advantages of a **single** XCNet over a classically learned CNet:

- ▶ extremely fast to learn
- ▶ only slightly less accurate as density estimators
- ▶ almost as good at generating samples
- ▶ less prone to overfitting

When plugged into **ensembles** they outperform state-of-the-art density estimators in a fraction of the time!

# XCNets II

Why a single XCNet is *not much worse* than a CNet?

Because of the mixture of experts interpretations, a path  $p = p_{(1)}p_{(2)} \cdots p_{(k)}$  connects the root to a single leaf  $\mathcal{M}_l$  after observing  $x_1x_2 \cdots x_k$ ,  $p_l(\mathbf{x}_{|\text{sc}(\mathcal{C}) \setminus \text{sc}(\mathcal{M}_l)})$  is decomposed by the chain rule across path  $p$ .

→ shuffling  $X_{p_{(0)}}, \dots, X_{p_{(k)}}$  does not influence inference!

Focus on learning **accurate local experts** (leaves) than the gating function!

Sampling is less affected:




CNet



Train



XCNet





Train

# XCNets III

Single C Nets learned with LearnCNet are prone to overfitting.

***Randomizing the gating function*** in XC Nets alleviate this issue.



Learning curves of C Nets and XC Nets (negative log-likelihoods) on KddCup2k and EachMovie The latter overfits much later than the former.

Moreover, it helps differentiating the local leaf experts



# Ensembles of XC Nets

Ensembles of XC Nets do not require to additionally diversify each components (no bootstrapping required).

Learning up to **500 components** of XC Nets is still **faster than learning 40** of other variants:

**CNet<sub>bag</sub>** bagging entCSN [Rahman and Gogate 2016]

**CNet<sub>boost</sub>** boosting entCSN [Rahman and Gogate 2016]

**dCSN<sup>k</sup>** bagging dCSN [Di Mauro, Vergari, and Basile 2015;  
Di Mauro, Vergari, and Esposito 2015]

# Experiments

We validate the following research questions:

- (Q1) how does a single XCNet **compare** to the optimal one learned by dCSN?
- (Q2) how **accurate** are ensembles of XC Nets compared state-of-the-art density estimators?
- (Q3) how **much time** do actually XC Nets save in practice?

We employ the **20 standard benchmark datasets** for density estimation [*Haaren and Davis 2012; Vergari, Mauro, and Esposito 2015*]

We compare to single and ensembles of C Nets plus other state-of-the-art TPMs like Sum-Product Networks (ID-SPN) and Markov Networks learned with ACs (ACMN) and even untractable Bayesian Networks (BN).

# (Q1) Single Model Comparison

| dataset    | entCNet       | dCSN           | XCNet              | dCSN <sub>PoB</sub> | XCNet <sub>PoB</sub> |
|------------|---------------|----------------|--------------------|---------------------|----------------------|
| NLTCS      | -6.06         | <b>-6.03</b>   | $6.06 \pm 0.01$    | -6.09               | $-6.17 \pm 0.05$     |
| MSNBC      | <b>-6.05</b>  | <b>-6.05</b>   | $-6.09 \pm 0.02$   | -6.05               | $-6.18 \pm 0.03$     |
| KDDCup2k   | -             | <b>-2.18</b>   | $-2.19 \pm 0.01$   | -2.19               | $-2.21 \pm 0.01$     |
| Plants     | <b>-13.25</b> | <b>-13.25</b>  | $-13.43 \pm 0.07$  | -14.89              | $-15.66 \pm 0.22$    |
| Audio      | <b>-42.05</b> | -42.10         | $-42.66 \pm 0.14$  | -42.95              | $-44.02 \pm 0.22$    |
| Jester     | -55.56        | <b>-55.40</b>  | $-56.10 \pm 0.19$  | -56.23              | $-57.39 \pm 0.15$    |
| Netflix    | <b>-58.71</b> | <b>-58.71</b>  | $-59.21 \pm 0.06$  | -60.20              | $-61.40 \pm 0.25$    |
| Accidents  | -30.69        | <b>-29.84</b>  | $-31.58 \pm 0.24$  | -36.24              | $-40.22 \pm 0.46$    |
| Retail     | <b>-10.94</b> | -11.24         | $-11.44 \pm 0.09$  | -11.06              | $-11.19 \pm 0.04$    |
| Pumsb-star | -24.42        | <b>-23.91</b>  | $-25.55 \pm 0.34$  | -32.11              | $-39.91 \pm 2.48$    |
| DNA        | -87.59        | <b>-87.31</b>  | $-87.67 \pm 0.00$  | -98.83              | $-99.84 \pm 0.05$    |
| Kosarek    | <b>-11.04</b> | -11.20         | $-11.70 \pm 0.13$  | -11.38              | $-11.80 \pm 0.07$    |
| MSWeb      | <b>-10.07</b> | -10.10         | $-10.47 \pm 0.10$  | -10.19              | $-10.43 \pm 0.07$    |
| Book       | <b>-37.35</b> | -38.93         | $-42.36 \pm 0.28$  | -38.21              | $-39.47 \pm 0.33$    |
| EachMovie  | -58.37        | <b>-58.06</b>  | $-60.71 \pm 0.89$  | -59.70              | $-62.58 \pm 0.38$    |
| WebKB      | -162.17       | <b>-161.92</b> | $-167.45 \pm 1.59$ | -168.7              | $-174.78 \pm 0.81$   |
| Reuters-52 | <b>-88.55</b> | -88.65         | $-99.52 \pm 1.93$  | -90.51              | $-100.25 \pm 0.57$   |
| 20NewsG    | -             | <b>-161.72</b> | $-172.6 \pm 1.40$  | -162.25             | $-167.39 \pm 0.74$   |
| BBC        | -263.08       | <b>-261.79</b> | $-261.79 \pm 0.00$ | -264.56             | $-274.83 \pm 1.15$   |
| Ad         | -16.92        | <b>-16.34</b>  | $-18.70 \pm 1.44$  | -36.44              | $-36.94 \pm 1.41$    |

Table 1.

Average test log-likelihoods for entCNet, dCSN, XCNet and their PoB variants dCSN<sub>PoB</sub> and XCNet<sub>PoB</sub>. For randomized models, mean and standard deviation over 10 runs are reported).

## (Q2) Ensemble Model Comparison

new *state-of-the-art* scores on 11/20 datasets...

| dataset    | CNet <sub>40</sub> <sub>bag</sub> | CNet <sub>40</sub> <sub>boost</sub> | dCSN <sub>40</sub> | XCNet <sub>40</sub> <sub>p08</sub> | XCNet <sub>40</sub> | XCNet <sub>500</sub> | ID-SPN        | ACMN          | WM            |
|------------|-----------------------------------|-------------------------------------|--------------------|------------------------------------|---------------------|----------------------|---------------|---------------|---------------|
| NLTCS      | <b>-6.00</b>                      | -6.01                               | <b>-6.00</b>       | -6.01                              | <b>-6.00</b>        | <b>-5.99</b>         | -6.02         | -6.00         | -6.02         |
| MSNBC      | -6.08                             | -6.15                               | <b>-6.05</b>       | -6.11                              | -6.06               | -6.06                | <b>-6.04</b>  | <b>-6.04</b>  | <b>-6.04</b>  |
| KDDCup2k   | -2.14                             | -2.15                               | -2.15              | <b>-2.13</b>                       | <b>-2.13</b>        | <b>-2.13</b>         | <b>-2.13</b>  | -2.17         | -2.16         |
| Plants     | -12.32                            | -12.67                              | -12.59             | -13.09                             | <b>-11.99</b>       | <b>-11.84</b>        | -12.54        | -12.80        | -12.65        |
| Audio      | -40.09                            | -39.84                              | -40.19             | -40.30                             | <b>-39.77</b>       | <b>-39.39</b>        | -39.79        | -40.32        | -40.50        |
| Jester     | -52.88                            | -52.82                              | -52.99             | -53.64                             | <b>-52.65</b>       | <b>-52.21</b>        | -52.86        | -53.31        | -53.85        |
| Netflix    | -56.55                            | -56.44                              | -56.69             | -57.64                             | <b>-56.38</b>       | <b>-55.93</b>        | -56.36        | -57.22        | -57.03        |
| Accidents  | -29.88                            | -29.45                              | <b>-29.27</b>      | -36.92                             | -29.31              | -29.10               | -26.98        | -27.11        | <b>-26.32</b> |
| Retail     | -10.84                            | <b>-10.81</b>                       | -11.17             | -10.88                             | -10.93              | -10.91               | -10.85        | -10.88        | -10.87        |
| Pumsb-star | -23.98                            | -23.46                              | -23.78             | -32.91                             | <b>-23.44</b>       | -23.31               | -22.41        | -23.55        | <b>-21.72</b> |
| DNA        | <b>-81.07</b>                     | -85.67                              | -85.95             | -98.28                             | -84.96              | -84.17               | -81.21        | <b>-80.03</b> | -80.65        |
| Kosarek    | -10.74                            | <b>-10.60</b>                       | -10.97             | -10.91                             | -10.72              | -10.66               | <b>-10.60</b> | -10.84        | -10.83        |
| MSWeb      | -9.77                             | -9.74                               | -9.93              | -9.83                              | <b>-9.66</b>        | <b>-9.62</b>         | -9.73         | -9.77         | -9.70         |
| Book       | -35.55                            | <b>-34.46</b>                       | -37.38             | -34.77                             | -36.35              | -35.45               | <b>-34.14</b> | -35.56        | -36.41        |
| EachMovie  | -53.00                            | <b>-51.53</b>                       | -54.14             | -51.66                             | -51.72              | <b>-50.34</b>        | -51.51        | -55.80        | -54.37        |
| WebKB      | -153.12                           | <b>-152.53</b>                      | -155.47            | -155.83                            | -153.01             | <b>-149.20</b>       | -151.84       | -159.13       | -157.43       |
| Reuters-52 | -83.71                            | <b>-83.69</b>                       | -86.19             | -85.16                             | -84.05              | <b>-81.87</b>        | -83.35        | -90.23        | -87.55        |
| 20NewsG    | -156.09                           | -153.12                             | -156.46            | <b>-152.21</b>                     | -153.89             | <b>-151.02</b>       | -151.47       | -161.13       | -158.95       |
| BBC        | <b>-237.42</b>                    | -247.01                             | -248.84            | -251.31                            | -238.47             | <b>-229.21</b>       | -248.93       | -257.10       | -257.86       |
| Ad         | -15.28                            | -14.36                              | -15.55             | -26.25                             | <b>-14.20</b>       | <b>-14.00</b>        | -19.05        | -16.53        | -18.35        |

## (Q3) Learning Times

...in a ***fraction of the time*** required by other competitors

| dataset    | dCSN  | XCNet | dCSN <sub>PoB</sub> | XCNet <sub>PoB</sub> | dCSN <sup>40</sup> | XCNet <sub>PoB</sub> <sup>40</sup> | XCNet <sup>40</sup> | XCNet <sup>500</sup> | ID-SPN |
|------------|-------|-------|---------------------|----------------------|--------------------|------------------------------------|---------------------|----------------------|--------|
| NLTCS      | 0     | 0.2   | 0.1                 | 0.01                 | 10                 | 0.2                                | 0.01                | 3                    | 310    |
| MSNBC      | 12    | 0.3   | 0.7                 | 0.01                 | 499                | 13.1                               | 13                  | 155                  | 46266  |
| KDDCup2k   | 112   | 0.5   | 12.0                | 0.32                 | 4126               | 21.2                               | 16                  | 247                  | 32067  |
| Plants     | 15    | 0.3   | 45.5                | 0.22                 | 325                | 1.0                                | 6                   | 77                   | 18833  |
| Audio      | 58    | 0.3   | 74.8                | 0.48                 | 980                | 0.8                                | 6                   | 136                  | 21009  |
| Jester     | 50    | 0.2   | 95.6                | 0.26                 | 989                | 0.3                                | 4                   | 83                   | 10412  |
| Netflix    | 75    | 0.2   | 2.8                 | 0.02                 | 1546               | 0.4                                | 9                   | 118                  | 30294  |
| Accidents  | 54    | 0.2   | 153.7               | 0.04                 | 996                | 0.7                                | 11                  | 138                  | 15472  |
| Retail     | 263   | 0.8   | 5.8                 | 0.01                 | 3780               | 3.2                                | 13                  | 164                  | 4041   |
| Pumsb-star | 118   | 0.6   | 26.2                | 0.02                 | 2260               | 0.8                                | 23                  | 290                  | 20952  |
| DNA        | 30    | 0.1   | 4.4                 | 0.01                 | 224                | 0.06                               | 3                   | 40                   | 3040   |
| Kosarek    | 588   | 2.4   | 41.2                | 0.01                 | 10033              | 10.8                               | 43                  | 524                  | 17799  |
| MSWeb      | 1215  | 7.2   | 7.4                 | 0.01                 | 17123              | 13.2                               | 129                 | 1592                 | 19682  |
| Book       | 9235  | 9.7   | 113.0               | 0.04                 | 155634             | 1.9                                | 316                 | 3476                 | 61248  |
| EachMovie  | 1297  | 7.1   | 4.7                 | 0.01                 | 16962              | 1.1                                | 127                 | 2601                 | 118782 |
| WebKB      | 4997  | 11.0  | 238.0               | 0.03                 | 18875              | 0.9                                | 190                 | 2237                 | 45451  |
| Reuters-52 | 9947  | 39.3  | 24.3                | 0.05                 | 65498              | 2.7                                | 414                 | 8423                 | 70863  |
| 20NewsG    | 16866 | 51.3  | 40.7                | 0.01                 | 153908             | 4.4                                | 506                 | 9883                 | 163256 |
| BBC        | 21381 | 8.4   | 7.3                 | 0.02                 | 69572              | 0.4                                | 256                 | 4251                 | 61471  |
| Ad         | 5212  | 116.5 | 134.0               | 0.08                 | 75694              | 4.2                                | 2403                | 30538                | 87522  |

Table 3. Times (in seconds) taken to learn the best models on each dataset for dCSN, XCNet, dCSN<sub>PoB</sub>, XCNet<sub>PoB</sub>, their ensembles and ID-SPN

# Conclusions

Due to their simplicity to implement, *fast learning times*, and *accurate inference performances*, XCNets set the **new baseline to compare against** for density estimation with TPMs (and not!).

## Future works

Exploiting the mixture of experts interpretation to devise more expressive gating functions to still perform exact and fast inference

## Code

<https://github.com/nicoladimauro/cnet>

## Paper

<http://www.di.uniba.it/~ndm/pubs/ndm17ecml.pdf>

# References

- ⊕ Di Mauro, Nicola, Antonio Vergari, and Teresa M.A. Basile (2015). "Learning Bayesian Random Cutset Forests". In: *Proceedings of ISMIS*. Springer, pp. 122–132.
- ⊕ Di Mauro, Nicola, Antonio Vergari, and Floriana Esposito (2015). "Learning Accurate Cutset Networks by Exploiting Decomposability". In: *Proceedings of AIXIA*. Springer, pp. 221–232.
- ⊕ Haaren, Jan Van and Jesse Davis (2012). "Markov Network Structure Learning: A Randomized Feature Generation Approach". In: *Proceedings of the 26th Conference on Artificial Intelligence*. AAAI Press.
- ⊕ Meilä, Marina and Michael I. Jordan (2000). "Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1, pp. 1–48.
- ⊕ Rahman, Tahrira and Vibhav Gogate (2016). "Learning Ensembles of Cutset Networks". In: *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*. AAAI'16. Phoenix, Arizona: AAAI Press, pp. 3301–3307. URL: <http://dl.acm.org/citation.cfm?id=3016100.3016365>.
- ⊕ Rahman, Tahrira, Prasanna Kothalkar, and Vibhav Gogate (2014). "Cutset Networks: A Simple, Tractable, and Scalable Approach for Improving the Accuracy of Chow-Liu Trees". In: *Machine Learning and Knowledge Discovery in Databases*. Vol. 8725. LNCS. Springer, pp. 630–645.
- ⊕ Roth, Dan (1996). "On the hardness of approximate reasoning". In: *Artificial Intelligence* 82.1–2, pp. 273–302.
- ⊕ Vergari, Antonio, Nicola Di Mauro, and Floriana Esposito (2015). "Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning". In: *ECML-PKDD 2015*.

## Discuss